

Hedge Funds on a Swamp: Analyzing Patterns, Vulnerabilities, and Defense Measures in Blockchain Bridges [Experiment, Analysis and Benchmark]

Poupak Azad
University of Manitoba
Canada
azad@myumanitoba.ca

Jiahua Xu
University College London
United Kingdom
jiahua.xu@ucl.ac.uk

Yebo Feng
Nanyang Technological University
Singapore
yebo.feng@ntu.edu.sg

Preston Strowbridge
University of Central Florida
USA
pr438045@ucf.edu

Cuneyt Gurcan Akcora
University of Central Florida
USA
cuneyt.akcora@ucf.edu

ABSTRACT

Blockchain bridges have become essential infrastructure for enabling interoperability across different blockchain networks, with more than \$24B monthly bridge transaction volume. However, their growing adoption has been accompanied by a disproportionate rise in security breaches, making them the single largest source of financial loss in Web3. For cross-chain ecosystems to be robust and sustainable, it is essential to understand and address these vulnerabilities. In this study, we present a comprehensive systematization of blockchain bridge design and security. We define three bridge security priors, formalize the architectural structure of 13 prominent bridges, and identify 23 attack vectors grounded in real-world blockchain exploits. Using this foundation, we evaluate 43 representative attack scenarios and introduce a layered threat model that captures security failures across source chain, off-chain, and destination chain components.

Our analysis at the static code and transaction network levels reveals recurring design flaws, particularly in access control, validator trust assumptions, and verification logic, and identifies key patterns in adversarial behavior based on transaction-level traces. To support future development, we propose a decision framework for bridge architecture design, along with defense mechanisms such as layered validation and circuit breakers. This work provides a data-driven foundation for evaluating bridge security and lays the groundwork for standardizing resilient cross-chain infrastructure.

1 INTRODUCTION

Blockchain bridges have become essential infrastructure in the blockchain ecosystem, enabling interoperability between otherwise isolated networks. In its most basic asset-based bridge type, a bridge locks or burns assets on a source chain and mints or releases corresponding assets on a destination chain. This coordination preserves total supply across networks while facilitating asset mobility and composability. Bridges play a foundational role in enabling cross-chain decentralized applications, unifying fragmented liquidity pools, and allowing users to leverage features across heterogeneous platforms. For instance, monthly bridge transaction volume has exceeded \$24B [1], underscoring their systemic significance.

However, this growing importance has been matched by an alarming trend: bridges represent the single largest source of financial loss in Web3 security breaches. As of mid 2025, 13 out of 39 bridges on l2beat.com are already labelled as insecure.¹ Of the top six bridges in total-value-locked as ranked by [2] in 2022, three (*Multichain* [3], *Ronin* [4], and *Rainbow* [5]) have already been hacked for more than \$750M. A critical vulnerability in the sixth, *Polygon's Plasma Bridge* [6], could have exposed \$850M; it was patched just in time after a white-hat disclosure that earned a \$2M bounty.

Vulnerabilities range from smart contract bugs and improper verification logic to compromised multisig keys and failure-prone trusted validators. The result is a persistent threat model with catastrophic consequences: stealthy attacks, rapid fund drainage, and minimal recourse for victims. The threat of large-scale, unpredictable failure continues to undermine user trust and impede the adoption of cross-chain systems. With billions of USD locked in bridges, their current security risks make them resemble hedge funds operating atop a swamp; highly valuable, yet dangerously unstable and exposed to unpredictable attacks.

Despite a growing body of research and security reviews, there is still no unified framework to evaluate bridge vulnerabilities. Existing literature tends to focus on isolated case studies [7, 8], post-mortem hack reports [9], or abstract protocol taxonomies [9, 10]. This fragmentation limits our ability to draw general conclusions, compare implementations, or develop preventive standards.

In this work, we take a data interoperability perspective to present a comprehensive security and privacy systematization of blockchain bridges, grounded in both formal modeling and empirical analysis. We argue that bridge security is fundamentally a data transfer security problem, concerned with state consistency, transaction ordering, and trust-minimized replication across chains. We believe that it must be addressed with the same rigor as traditional interoperability standards.

We approach the study of bridge security through three orthogonal lenses:

¹The 13 bridges marked as vulnerable (indicated with a red cross and / or brown shield badge on the website) due to unverified contracts and past hacks are: *Aptos* (LayerZero), *LayerZero v2 OFTs*, *Multichain*, *Connex*, *Omnichain* (LayerZero), *Hyperlane*, *Nexus*, *Socket*, *Chainport*, *Allbridge*, *StarGate* (LayerZero), *Symbiosis*, *Everclear*, and *Orbit Bridge*.

First, we formalize the operational layers of bridge systems, spanning the source chain, off-chain intermediaries, and destination chain components. This layered model allows us to define three core security priors, such as cross-chain causality, and to reason about how different bridge designs instantiate or violate them.

Second, we perform a large-scale static analysis of bridge smart contracts deployed on Ethereum [11]. Using program analysis and custom metrics, we map the usage of role-based access control, defensive programming patterns, error handling, and call structures across 13 major bridge protocols. Next, we extract bridge-related transaction networks from on-chain data and apply graph machine learning to analyze usage and attacker behavior in real-world incidents. Our analysis spans 43 major attacks and reveals patterns in exploit phases, laundering strategies, and response delays.

Most importantly, we aim to answer four research questions to evaluate proactive goals: **RQ1** (Prevention): To what extent are bridge attacks preventable? Which defense techniques, such as static code analysis, formal verification, or trusted execution environments, are effective? Are some bridge types inherently more secure? **RQ2** (Detection): If prevention is not guaranteed, can attacks be detected in real time? What mechanisms (e.g., anomaly detection, on-chain monitoring) support timely identification? **RQ3** (Mitigation): Once an attack is detected, what mechanisms exist to mitigate damage and minimize asset loss? How do mitigation strategies differ across bridge types? **RQ4** (Outlook): What is the long-term outlook for bridge security? What foundational steps are needed to design formally verifiable or provably secure bridge protocols?

By unifying theoretical foundations with large-scale empirical analysis, this work offers the first comprehensive, data-driven systematization of blockchain bridge security. Our findings serve as a baseline for future research, standardization efforts, and the secure evolution of bridge protocols.

Our contributions are as follows:

- We formalize a layered model of blockchain bridge architectures and define core security priors relevant to cross-chain behavior.
- We develop a vulnerability taxonomy and introduce a formal notion of bridge attack surfaces based on trust assumptions and implementation details.
- We perform the first large-scale static analysis of deployed bridge contracts, quantifying security patterns across access controls, call structures, and guard mechanisms.
- We extract and analyze transaction-level behavior from bridge exploit incidents, identifying behavioral patterns across phases of use, compromise, and fund laundering.
- We provide security benchmarks and design recommendations to guide future bridge development, formal analysis, and regulatory evaluation.

2 RELATED WORK

Recent literature reflects an active effort to reconcile decentralization, trust minimization, and performance in blockchain bridges. We summarize case studies and surveys here and refer the reader to the supplementary material (available at <https://fdatalab.github.io/BlockchainBridgeAnalysis/>) for additional details on interoperability and benchmark studies.

Transaction analysis for bridges can identify operational patterns, including usage behaviors, anomalous activity, and indicators of compromise. Huang et al. [12] present an in-depth empirical analysis of Stargate, a prominent Layer-0 bridge with one of the highest total value locked (TVL) figures in the ecosystem. Using on-chain data from six EVM-compatible blockchains (Ethereum, Polygon [13], BSC [14], Avalanche [15], Arbitrum [16], Optimism [17]), they examine Stargate’s transaction volume, user adoption, and operational patterns. This case study provides a data-driven benchmark for analyzing cross-chain liquidity transfers in practice. However, the study does not identify bridge priors nor carry out a static code analysis.

Subramanian et al. [18] benchmark the performance of blockchain bridge aggregators: services that route transfers across multiple bridge protocols for optimal speed or cost. They develop a framework to test popular aggregators (e.g., LLFI, Socket, deBridge) by executing hundreds of cross-chain swaps and recording metrics such as fees, slippage, and latency. Their findings quantify differences in cost-efficiency and reliability, providing performance benchmarks for user-centric interoperability services. Our work differs in its broader focus and the development of a taxonomy to offer a holistic view of bridges.

Augusto et al. [19] introduce *XChainDataGen*, a framework for generating large-scale bridge transaction datasets. Using this tool, they collect approximately 35 GB of data from five major bridges deployed on 11 blockchains during the second half of 2024, extracting over 11.2 million bridge transactions that moved over \$28B in token value. They compare protocols in terms of security, cost, and performance, contrasting, for example, full source-chain finality with “soft” finality, alongside fee models and emerging paradigms such as cross-chain intents. However, the study focuses on financial aspects and does not study bridge vulnerabilities that hinder bridge adoption in finance.

Recent surveys and taxonomies have focused on identifying systemic vulnerabilities in bridge designs. Li et al. [10] provide a comprehensive review of blockchain bridges, classifying architectures (e.g., lock-mint vs. notary schemes) and documenting common vulnerabilities such as smart contract flaws, centralization risks, liquidity issues, and oracle manipulation. While the study presents a broad taxonomy and useful design insights, it does not analyze transaction data or investigate real-world exploits. In a recent work, Augusto et al. [20] analyze bridge exploits; however, they do not conduct a direct analysis of what happens to stolen funds after bridge hacks, nor carry out static code analysis themselves. Instead, they rely heavily on systematic literature review, audit reports, bug bounty disclosures, and gray literature to collect information about vulnerabilities and attacks. Our analysis covers both the code and transaction analysis.

In contrast, Belenkov et al. [9] present a Systematization of Knowledge focused on major bridge hacks, including the \$600M Axie *Ronin* exploit. They categorize failure modes such as compromised keys, multisig errors, and validator logic flaws, and propose best practices for prevention. However, the analysis is limited to static perspectives and excludes transactional behavior. Our work complements and extends these efforts by combining both static and transaction analysis perspectives and grounds them in a unified formal model.

Most security and forensic analyses of bridge hacks have been retrospective machine learning studies of transaction networks, typically conducted on a per-case basis and without developing generalizable vulnerability taxonomies or formalizing bridge-specific assumptions. For example, Augusto et al. [7] propose *XChainWatcher*, a Datalog-based monitoring system that detects anomalous token flows across bridges in real time. While it successfully reconstructs the *Ronin* and *Nomad* [21] attacks, its scope is limited to two case studies and does not address broader questions about interoperability security. Similarly, Lin et al. [22] develop *ABCTracer*, an automated framework for linking bridge transaction legs. Combining event log mining with machine learning inference, *ABCTracer* achieves 91.75% F1-score in identifying transaction pairs across 12 DeFi bridges. While valuable for forensics, the work does not define a vulnerability taxonomy nor investigate design assumptions underpinning bridge protocols.

Finally, Wu et al. [8] focus on identifying and classifying bridge-specific attacks. Using a dataset of 49 bridge exploits totaling nearly \$4.3B in losses, they propose *BridgeGuard*, a graph-based detection system that models bridge transactions and flags anomalous behavior. Evaluated on 203 known attacks and 40,000 benign transactions, *BridgeGuard* outperforms prior methods in recall and detection precision. However, their work does not incorporate static analysis of contract logic or propose a general framework for understanding bridge security.

While these studies contribute valuable empirical and forensic insights, they tend to focus narrowly on individual case studies, static audits, or retrospective detection. In contrast, our work presents the first unified and multi-dimensional study that combines a formal model of bridge semantics, a systematic static analysis of deployed bridge contracts, and a large-scale transaction analysis across multiple chains.

3 BRIDGE MODELING AND FORMALIZATION

A blockchain b is an immutable ledger where transactions are appended chronologically. It comprises a set of peer-to-peer network nodes N_p , a set of address nodes N_a , a chronological history of transactions $\mathbb{H}$ between nodes in N_a , and consensus rules R that govern transaction creation, where $N_p \cap N_a \neq \emptyset$ and $b = (N_p, N_a, \mathbb{H}, R)$.

Asset Types. A clear distinction exists between coins, tokens, and wrapped assets. Coins are native digital currencies intrinsic to their own blockchains, such as Bitcoin (BTC) on the Bitcoin network [23], and are essential for paying transaction fees and participating in consensus mechanisms. A smart contract-based token θ represents a type of blockchain currency, with its state defined by the chronological history of transactions $\mathbb{H}_\theta$ in which θ has been involved. All blockchains issue native coins (e.g., Ether on Ethereum), but only some blockchains allow users to issue smart contract-based tokens (e.g., Storj on Ethereum).

Wrapped tokens are a subset of tokens designed to represent tokens from one blockchain on another, facilitating interoperability across disparate blockchain networks. For instance, Wrapped Bitcoin (WBTC) is issued on Ethereum, enabling BTC to be utilized within Ethereum’s decentralized finance ecosystem. Despite

enhancing bridge functionality, wrapped tokens inherit the limitations of standard tokens: they cannot be used to pay transaction fees on the host blockchain (i.e., Ethereum for WBTC).

A transaction $tx \in \mathbb{H}$ on b transfers a value v of a token θ from an address $a_1 \in N_a$ to an address $a_2 \in N_a$, and is associated with a timestamp t_{tx} indicating when the transaction occurred: $tx = (\theta, v, a_1, a_2, t_{tx})$. If the transaction is understood, we will simplify t_{tx} to t . Each transaction is assumed to be valid under the consensus rules R of b .

Blockchain Layers. Blockchain systems are structured into hierarchical layers to manage scalability, functionality, and specialization.

- **Layer 1 (L1):** Base layer of a blockchain network, encompassing the core protocol responsible for transaction validation, consensus mechanisms, and data storage. L1 blockchains operate independently and are the foundation upon which other layers are built. Examples include Bitcoin, Ethereum, and Solana [24].
- **Layer 2 (L2):** Built atop L1 blockchains, L2 protocols aim to enhance scalability and transaction throughput without altering the base protocol. They achieve this by processing transactions off-chain or in parallel, subsequently settling them on the L1. Notable L2 implementations include Optimism and Arbitrum on Ethereum.
- **Layer 3 (L3):** L3 represents an emerging concept focusing on application-specific functionalities. These are protocols or networks constructed on L2 solutions, offering tailored environments for decentralized applications (dApps). L3s aim to provide enhanced scalability, interoperability, and customization, facilitating complex applications like decentralized finance platforms and gaming ecosystems. Examples of L3 projects include Orbs Network [25] and XAI Games [26] on Arbitrum.

3.1 Blockchain Bridges

A blockchain bridge facilitates the transfer of tokens across distinct blockchain domains. While many bridges connect L1 blockchains, others operate across layers, such as the *Arbitrum Canonical Bridge* between Ethereum (L1) and Arbitrum (L2). Bridges that involve L3 layers are an emerging area with early examples such as *Arbitrum Orbit*, *zkSync Hyperchains* [27], and *Orbs Network*. Despite their growing importance, blockchain bridges lack a consistent formalization in the literature, which we aim to define as follows.

We consider two blockchain domains (i.e., chains or protocols), $b_1 \in \{L1, L2, L3\}$ and $b_2 \in \{L1, L2, L3\}$, which may differ in node sets, transaction histories, and consensus rules. We define a bridge between the domains b_1 and b_2 through an implementation mechanism $\mathcal{I}$ that governs the operational logic of the bridge:

$$\mathbb{B}_{1 \leftrightarrow 2} = (\{b_1, b_2\}, \mathcal{I}) \quad (1)$$

We categorize bridges along two orthogonal axes: their *operational trust model* and their *functional type*. The trust model refers to how source domain activity is verified and includes *trusted*, *trust-minimized*, and *trustless* designs, formalized shortly in Section 3.5. The functional type captures how value is transferred and includes:

- **Asset Bridges:** In *asset bridges* (also known as burn-and-mint models), the bridge locks or burns an asset θ_1 on the source domain b_1 and creates a representative asset θ_2 on the destination

domain b_2 . These assets are not inherently equivalent; their linkage is established solely through the semantics of the bridge protocol. They include:

- *Externally-Verified Asset Bridges*, which depend on multisigs, notaries, or sidechains to verify lock events (e.g., *Wormhole* [28], *Ronin*, *Multichain*).
- *Rollup-Native Asset Bridges*, which leverage L1 consensus to verify L2 state transitions using fraud or validity proofs (e.g., Arbitrum, Optimism, Loopring [29]). These are considered the most trustless.

The settlement process on asset-based bridges may be slow and require specialized verification and challenge code. For example, optimistic rollups like Arbitrum or Optimism impose a 7-day challenge period on withdrawals.

- **Liquidity Networks:** In contrast to asset-based bridges, *liquidity network-based bridges*, such as Connex[30] or Across[31], avoid minting and instead fulfill user requests through liquidity providers (LPs) who maintain reserves on both domains. LPs are economically incentivized via transfer fees and, in some protocols, additional yield or reward mechanisms. This design enables faster settlement and circumvents the latency of on-chain verification and challenge periods.
- **Hybrid Bridges:** These support both functional models, often depending on the specific chain pair or asset type. For instance, *Multichain* combines canonical minting with liquidity provisioning.

A bridge transaction tx moves a value v of token θ_1 from an address $a_1 \in N_a$ on b_1 to an address $a_2 \in N_a$ on b_2 :

$$tx = (\theta_1, v, a_1, a_2, t_{tx}) \quad (2)$$

The transformation of a bridge transaction state χ under the bridge's operation is defined as:

$$\chi_0 \xrightarrow{\mathbb{B}_{1 \leftrightarrow 2}} \chi_t \quad (3)$$

where χ_0 and χ_t represent the initial and final states of the transaction across chains.

The implementation $\mathcal{I}$ of the bridge includes mechanisms on the source chain, off-chain components, and the destination chain. Additionally, the bridge relies on a node trust set $\mathbb{T}$, comprising entities responsible for validating and securing bridge transactions. The composition of $\mathbb{T}$ is determined by the specific implementation $\mathcal{I}$ and the security assumptions of $\{b_1, b_2\}$.

3.2 Bridge Mechanism

Consider that Alice wants to move some assets from the source chain b_1 to the destination chain b_2 . For simplicity in exposition, we will assume that at a given time, the bridge is used by one user only and for bridging one token only. The token Alice will be moving is θ_1 , and the representation of the token on b_2 is θ_2 . We track the movement of value across the bridge $\mathbb{B}_{1 \leftrightarrow 2}$ as a function of time and fee structure. The state of the bridge transaction χ is described by four accounts: Alice's addresses a_1 on b_1 and a_2 on b_2 along with the bridge addresses c_1 on b_1 and c_2 on b_2 . The state of the bridge transaction is represented as $\chi = (a_1, c_1, a_2, c_2)$.

We track the mechanisms in three stages: source chain mechanism, off-chain mechanism, and destination chain mechanism.

3.2.1 Source Chain Mechanism. Alice initiates the transfer on b_1 . Initially, at $t = 0$, Alice holds v_1 units of the token θ_1 at her address a_1 , with a total value $v_1 \cdot \text{price}(\theta_1, t)$, where v_1 is the amount of the token, and $\text{price}(\theta_1, t)$ represents the price of the token θ_1 in a fiat currency such as USD at time t . We will use $a_x \mapsto v$ to show that the balance of address a_x is v . Hence, the initial state of balances is $\chi_0 = (a_1 \mapsto v_1, c_1 \mapsto 0, a_2 \mapsto 0, c_2 \mapsto 0)$.

If Alice wants to transfer an amount $v_x \leq v_1$ to her address a_2 on the other blockchain, she initiates the transaction $tx_{b_1} = (\theta_1, v_x, a_1, c_1, t_{tx_{b_1}})$ on b_1 .

She also pays a bridge fee F_{forward} , which consists of the transaction processing fee f_1 on b_1 and f_2 on b_2 , as well as the bridge operation fee (f^*). Hence the total fee $F_{\text{forward}} = f_1 + f_2 + f^*$. For simplicity, we assume that the bridge operation fee is charged on the initiating blockchain b_1 . The sent amount v_x and fees are deducted from the initial balance of $a_1 \mapsto v_1$, hence $a_1 \mapsto (v_1 - v_x - f_1 - f^*)$. This leaves only f_2 to be paid on b_2 .

After b_1 's block confirmation duration d_{b_1} , the bridge smart contract on b_1 acknowledges tx_{b_1} and locks ($c_1 \mapsto v_x$) or burns ($c_1 \mapsto 0$) the received assets depending on the implementation $\mathcal{I}$. We follow the locked asset scenario, and the state is updated as $\chi_{t_{tx_{b_1}}} = (a_1 \mapsto (v_1 - v_x - f_1 - f^*), c_1 \mapsto v_x, a_2 \mapsto 0, c_2 \mapsto 0)$. The blockchain state changes as $b_1 = (N_p, N_a, \mathbb{H} \cup \{tx_{b_1}\}, R)$.

3.2.2 Off-chain Mechanism. An off-chain communicator observes transactions to bridge address c_1 , specifically transactions of the form: $tx_{B_1} = (\theta, v, a, c, t)$. Upon detection, it instructs c_2 to mint $v_x - f_2$ worth of θ_2 . The off-chain mechanism takes time d_{off} to notice the first transaction and send a signal to the b_2 . Thus, at $t_{\text{off}} = t_{tx_{b_1}} + d_{\text{off}}$, the off-chain mechanism signals the start of the token transfer process on the second blockchain b_2 . Specifically, the mechanism signals the smart contract on b_2 accordingly. There are multiple mechanisms to create this signal, and we refer the reader to [9] for a comprehensive list. In a simple oracle-based solution, a network of oracles monitors transactions on blockchain b_1 . These relayers are either permissioned (managed by a specific entity) or decentralized (e.g., a set of staked validators). When a bridge transaction tx_{b_1} is observed, relayers submit proofs or messages (through transactions that push data to the blockchain) to a contract on b_2 to trigger the minting or unlocking of assets. The node trust set $\mathbb{T}$ validates these processes based on the implementation $\mathcal{I}$.

3.2.3 Destination Chain Mechanism. The transfer completes when $v_x - f_2$ worth of θ_2 is minted on b_2 and sent to Alice's address a_2 : $tx_{b_2} = (\theta_2, v_2 = 0, a_2 \mapsto v_x - f_2, c_2 \mapsto 0, t_{\text{off}})$. After b_2 's block confirmation duration D_{b_2} , the complete bridge process ends at $t = t_{tx_{b_1}} + d_{\text{off}} + d_{b_2}$. The final state is: $\chi_t = (a_1 \mapsto (v - v_x - f_1 - f^*), c_1 \mapsto 0, a_2 \mapsto v_x - f_2, c_2 \mapsto 0)$.

The destination domain is $b_2 = (N_p, N_a, \mathbb{H} \cup \{t_{b_2}\}, R)$. The smart contract on b_2 handles the transaction and updates the blockchain state accordingly.

3.2.4 Reverse Process. The reverse transfer follows the same mechanism as the forward transfer but in the opposite direction. Alice initiates a transaction on b_2 to send v_x of θ_2 back to b_1 . The bridge smart contract on b_2 locks ($c_2 \mapsto v_x$) or burns ($c_2 \mapsto 0$) token θ_2 , and after confirmation, an off-chain mechanism signals b_1 to release

$(v_x - F_{\text{reverse}})$ of θ_1 to Alice's address a_1 . The final state transition mirrors the forward process, with updated fees and timestamps.

3.3 Formalization of Bridge Implementations

Blockchain bridges can be categorized based on their trust assumptions, yielding three classes: trustless, trusted, and trust-minimized. A trustless bridge has no external trusted entities, with $\mathbb{T}_{\text{trustless}} = \{\}$, relying exclusively on the security guarantees of the underlying blockchains. A trusted bridge introduces external dependencies, requiring $\mathbb{T}_{\text{trusted}} \neq \{\}$. A trust-minimized bridge is a subset of trusted bridges in which all trusted entities $\mathbb{T}_{\text{trustmin}}$ consist of deterministic, publicly auditable algorithms (for example, on-chain smart contracts, oracles, or relays). While trust-minimized bridges rely on external components, their correctness can be independently verified by users.

The bridge trust set is expressed as $\mathbb{T} = \mathbb{T}_{\text{src}} \cup \mathbb{T}_{\text{off}} \cup \mathbb{T}_{\text{dest}}$, where $\mathbb{T}_{\text{src}}$ captures trust assumptions on the source blockchain, $\mathbb{T}_{\text{off}}$ captures trust in off-chain mechanisms (such as relayers or validators), and $\mathbb{T}_{\text{dest}}$ captures trust in the destination blockchain. Two key metrics, $\text{Size}(\mathbb{T})$ and $\text{cost}(\mathbb{T})$, characterize the scale of the trusted set and its overall fee impact, respectively. These metrics influence the security, decentralization, and efficiency of a bridge.

3.3.1 Source Chain Implementation for Trust Set $\mathbb{T}_{\text{src}}$. The source chain component locks or burns assets before triggering $\mathbb{T}_{\text{off}}$. It can be realized via smart contracts (so $\mathbb{T}_{\text{src}} = \{\text{SCs}\}$), validator-controlled mechanisms ($\mathbb{T}_{\text{src}} = \{\}$), or hybrid approaches that combine both. Smart contract mechanisms require trusting the correctness of the contract logic, whereas validator-based methods rely on the blockchain's consensus, eliminating additional trust. Hybrid methods usually trade off fees or time for scalability and security. As shown in Table 1, these components appear on both source and destination chains, but since their role is more critical on the destination side, where bridged tokens are released, we omit their detailed explanation here and defer it to the destination chain discussion to avoid repetition.

3.3.2 Off-chain Implementation for Trust Set $\mathbb{T}_{\text{off}}$. Bridges typically use either a set of notaries (N) or a set of light clients (L) for off-chain processing.

Notaries. Notaries introduce additional trust in external parties ($\mathbb{T}_{\text{off}} \neq \{\}$) that monitor transactions and ensure bridge causality, with security depending on the number of notaries $\text{Size}(N)$ and their operational costs $\text{cost}(N)$. As $\text{Size}(N)$ and $\text{cost}(N)$ grow, security improves, but so do fees and transaction delays. The expected loss costs $\mathbb{E}(C)$, bridge fees F , and transaction time delay D follow $\mathbb{E}(C) \propto 1/\text{Size}(N)$ and $F, D \propto \text{Size}(N), \text{Cost}(N)$.




Figure 1: Light Client Implementation

Light Clients. Light clients rely on Merkle proofs (M) and perform verification on-chain via light client contracts (L). These systems do not depend on external entities, so the off-chain trust set is $\mathbb{T}_{\text{off}} = \{L, M\}$. These bridges are trust-minimized because no external entities are needed beyond the blockchain itself. Verification time t_{proof} affects expected loss costs $\mathbb{E}(C)$ and bridge cost F, D according to $\mathbb{E}(C) \propto 1/t_{\text{proof}}$, $F, D \propto t_{\text{proof}}$. Although frequent proof generation boosts security, it also increases fees. Light clients must run on each chain, reducing overall scalability. Figure 1 illustrates a standard light client implementation used in many trustless bridge designs.

Some bridges blend the two approaches. A security-optimized hybrid requires that either the light client or the notary set remain honest, so $\mathbb{T}_{\text{off}} = \{L, M\} \cap \{N\}$. Another hybrid uses light clients where available and defaults to notaries otherwise: $\mathbb{T}_{\text{off}} = \{L, M\} \vee \{N\}$.

Sidechains. A third option is using sidechains, which function as independent blockchains that can facilitate bridge transfers. Sidechains add a secondary blockchain with its own consensus R' . A bridge passing through a sidechain depends on $\mathbb{T}_{\text{off}} = (\{L, M\} \vee \{N\} \vee (\{L, M\} \cap \{N\})) \cup R'$.

In other words, instead of relying exclusively on a notary network or on-chain light clients to handle the bridging, a sidechain can be set up with built-in asset-transfer features. Bridge transactions then pass through that sidechain, inheriting its consensus and security properties.

While the idea of a sidechain can still incorporate notaries or light clients under the hood, the main distinction is that a sidechain is an entire blockchain in its own right, rather than a narrowly defined tool like a light client or a notary service. This extra layer can increase scalability (since the sidechain can process many transactions internally) but also introduces new trust assumptions. Specifically, the correctness and security of the sidechain's consensus mechanism. Figure 2 illustrates a representative sidechain-based bridge design.




Figure 2: Sidechain Implementation

3.3.3 Destination Chain Implementation $\mathbb{T}_{\text{dest}}$. The destination side enforces token minting, final validation, and recipient assignment. Possible methods include smart contracts ($\mathbb{T}_{\text{dest}} = \{\text{SCs}\}$), validator-control ($\mathbb{T}_{\text{dest}} = \{\}$), custodian-based ($\mathbb{T}_{\text{dest}} = \{C\}$), or a hybrid that integrates these techniques. The choice depends on constraints related to cost, security, and overall efficiency.

Smart Contracts. A smart contract can be deployed on the destination chain to handle final token transfers. After receiving a valid signal and data from $\mathbb{T}_{\text{off}}$, the contract locks, mints, or releases

Table 1: Cross-chain bridge classification for the highest volume bridges. ●: trusted, ◐: trust minimized, ○: trustless.

	Bridge	Source Chain	Destination Chain	Trust
Notaries	Nomad Bridge	Hybrid	Smart Contracts	●
	Allbridge Classic	Smart Contracts	Smart Contracts	●
	deBridge	Smart Contracts	Smart Contracts	●
	Multichain	Smart Contracts	Smart Contracts	●
	Wormhole Bridge	Smart Contracts	Smart Contracts	●
	Avalanche Bridge	Smart Contracts	Validator Control	●
	Ronin Bridge	Smart Contracts	Smart Contracts	●
	Wanchain Bridge	Smart Contracts	Smart Contracts	●
Light Client	BTC Relay	Validator Control	Smart Contracts	◐
	zkBridge	Validator Control	Validator Control	◐
	Rainbow Bridge	Validator Control	Validator Control	◐
	PeaceRelay	Smart Contracts	Smart Contracts	◐
Hybrid	Connex	Smart Contracts	Smart Contracts	●
	Optics Bridge	Smart Contracts	Smart Contracts	●
Sidechain	Cosmos IBC	Validator Control	Validator Control	○
	Gravity Bridge	Smart Contracts	Smart Contracts	◐
	zkRelay	Smart Contracts	Smart Contracts	◐
	Cactus	Smart Contracts	Smart Contracts	◐
	Celer cBridge	Smart Contracts	Smart Contracts	◐
	Orbit Bridge	Smart Contracts	Smart Contracts	◐
	Axelar	Smart Contracts	Smart Contracts	◐
	Chainswap	Smart Contracts	Smart Contracts	◐
	PolyBridge	Smart Contracts	Smart Contracts	◐
	pNetwork	Smart Contracts	Smart Contracts	◐
	Meter Passport	Smart Contracts	Smart Contracts	◐
	QANX Bridge	Smart Contracts	Smart Contracts	◐
	Binance Bridge	Smart Contracts	Smart Contracts	◐
	Horizon Bridge	Smart Contracts	Smart Contracts	◐
	Plasma Bridge	Smart Contracts	Smart Contracts	◐

tokens to the designated address. Because users must trust the correctness of the contract code, we have $\mathbb{T}_{dest} = \{SC\}$.

Validator Control. If the destination blockchain’s existing consensus participants validate the bridging process, no additional trusted entities are introduced. The bridge may still have predefined logic to identify the destination address, but final checks are performed by decentralized validators. Thus $\mathbb{T}_{dest} = \{\}$.

Hybrid. A hybrid model combines smart contracts with validator-based consensus to balance security, cost, and time. The exact composition of $\mathbb{T}_{dest}$ depends on how responsibilities are allocated between on-chain code and validator consensus.

3.4 Bridge Security Priors

We formalize three key security properties that bridges must satisfy to ensure token parity, transaction causality, and value preservation across blockchains. A fourth foundational requirement is the liveness of the bridge, but we assume that the underlying blockchains are live and capable of processing transactions in a timely manner, allowing us to focus on security properties specific to the bridge itself.

Bridge Peg. Blockchain domains b_1 and b_2 often have disparate implementations of tokens. For users to reliably use token θ_1 on b_1 as a representation of token θ_2 on b_2 , the bridge must establish that the asset amounts remain equivalent, i.e., $v_1 \equiv v_2$. This holds only if a peg exists between θ_1 and θ_2 , ensuring that one θ_1 can always be exchanged for one θ_2 . For simplicity, we ignore the total fee $F_{forward}$ that would be incurred to use the bridge, and state that

given token prices at time t , denoted as $price(\theta_1, t)$ and $price(\theta_2, t)$, the bridge must ensure

$$v_1 \cdot price(\theta_1, t) \equiv v_2 \cdot price(\theta_2, t), \quad \forall t. \quad (4)$$

To uphold this peg, the bridge must maintain both liveness and security. If the bridge experiences technical failure or an attack, equivalency in (4) may break down, leading to price divergence between the native and bridged tokens, which can disrupt user trust and capital efficiency.

Bridge Causality. A secure bridge must enforce causality, ensuring that no user extracts value that was never created or loses value that was not meant to be lost. Formally:

$$\begin{aligned} \forall tx_{B_2} = (\theta_2, \dots, c_2, t_2), \quad \exists! tx_{B_1} = (\theta_1, \dots, c_1, t_1) \\ \forall tx_{B_1} = (\theta_1, \dots, c_1, t_1), \quad \exists! tx_{B_2} = (\theta_2, \dots, c_2, t_2) \end{aligned} \quad (5)$$

such that $t_1 < t_2$.

This ensures a bijective mapping between transactions on b_1 and b_2 , preventing double-minting or loss of funds. Additionally, the bridge must preserve temporal ordering, enforcing $t_1 < t_2$ so that a transaction on b_1 must occur before its corresponding transaction on b_2 .

Bridge Consistency. The bridge must guarantee that tokens locked on b_1 remain inaccessible until the corresponding tokens on b_2 are either burned or locked. This consistency constraint ensures that the system does not create or destroy value arbitrarily.

$$v_2 \neq 0 \Rightarrow \neg \exists tx_{B_1}(\theta_2, v_1, c_1) \quad (6)$$

Equation (6) prevents a user from simultaneously withdrawing locked assets on b_1 and holding minted tokens on b_2 , which could result in infinite money creation.

$$v_2 > 0 \Rightarrow \exists! \theta_1 \text{ such that } c_1 \mapsto a_2. \quad (7)$$

Equation (7) strengthens this by enforcing a strict one-to-one correspondence between locked and minted tokens.

Value at Risk. Failures in bridge security can lead to asset loss with custodial, contract, and economic risks. In custodial risk, an externally controlled address or multi-sig wallet is compromised, and funds are permanently stolen. In contract risk, a smart contract-based bridge has a vulnerability (e.g., reentrancy attack) where attackers can extract tokens illicitly. In economic risk, a bridge relies on economic assumptions (e.g., optimistic rollups or bonded validators), but adversaries manipulate incentives to destabilize the system.

Losses may occur at i) the bridge contract level, where contract c_1 on b_1 and contract c_2 on b_2 impact all users of the bridge, or ii) at the user level, where losses occur for individual user addresses a_1 or a_2 . User-level losses (e.g., a_2 is not a valid address) are orthogonal to bridge security as long as they are not caused by flaws in bridge design.

A bridge’s trust model (custodian-based, validator-based, or smart contract-based) influences its risk profile and determines how funds are secured.

3.5 Formalization of Bridge Attack Surfaces

Current research on bridge attacks is fragmented, lacking a consistent definition of attack surfaces and vectors. We present a formalized model for analyzing bridge security.

DEFINITION 3.1 (ATTACK). *If (4), (5), (6), or (7) are violated due to malicious agents, the bridge is under attack.*

DEFINITION 3.2 (FAILURE). *If (4), (5), (6), or (7) are violated due to technical errors but without malicious intent, the bridge is experiencing failure.*

THEOREM 1. *If a bridge experiences an attack or failure, then (4) is always violated.*

Due to space limitations, the proof is given in the supplementary material.

DEFINITION 3.3 (ATTACK VECTOR V). *A vulnerability, pathway, or method that a malicious agent can exploit to launch an attack on a bridge $\mathbb{B}_{1 \leftrightarrow 2}$.*

DEFINITION 3.4 (ATTACK SURFACE Σ). *The sum of all attack vectors for a given $\mathbb{B}_{1 \leftrightarrow 2}$. Formally, an attack surface is defined as $\Sigma = \langle \mathbb{T}, \mathbb{I} \rangle$, where $\mathbb{T}$ represents trusted entities and $\mathbb{I}$ represents implementation details.*

An attack surface can be decomposed into sub-surfaces for finer-grained analysis. A sub-surface $\Sigma'_{\mathbb{B}_{1 \leftrightarrow 2}} \in \Sigma_{\mathbb{B}_{1 \leftrightarrow 2}}$ is defined as:

$$\Sigma'_{\mathbb{B}_{1 \leftrightarrow 2}} = \langle \mathbb{T}', \mathbb{I}' \rangle, \quad \text{where } \mathbb{T}' \in \mathbb{T}, \quad \mathbb{I}' \in \mathbb{I}.$$

Attack vectors in $\Sigma'_{\mathbb{B}_{1 \leftrightarrow 2}}$ use all trusted entities $\mathbb{T}'$ and span specific implementations $\mathbb{I}'$.

3.5.1 Examples of Attack Surfaces. In custodial key compromise ($\mathbb{T} = \{C\}$), a centralized custodian's private key is compromised, and all assets in the bridge are at risk. In the reentrancy attack on smart contract-based bridges ($\mathbb{T} = \{SCs\}$), an attacker exploits an improper withdrawal function to repeatedly drain funds. In validator collusion ($\mathbb{T} = \{V\}$), the validators of a PoS-based bridge collude; they approve fraudulent transactions and steal assets.

3.5.2 An Attack Surface-Damage Model

DEFINITION 3.5 (DAMAGE POTENTIAL/EFFORT RATIO). *The Damage Potential/Effort Ratio of an attack vector V , denoted $der(V)$, quantifies the feasibility of an attack and is defined as $der(V) = \frac{I(V)}{E(V)}$, where $I(V)$ represents the expected damage or impact of the attack vector and $E(V)$ represents the computational, economic, or procedural effort required to execute the attack.*

An attack vector V is viable if and only if $der(V) > 1$, indicating that the expected damage outweighs the effort required to execute the attack. Hence,

$$der(V) = \begin{cases} 1, & \text{if } der(V) > 1 \\ 0, & \text{otherwise} \end{cases}$$

We define the attack surface area of a bridge $\mathbb{B}_{1 \leftrightarrow 2}$, or a subset of it, containing n attack vectors as:

$$Area(\Sigma'_{\mathbb{B}_{1 \leftrightarrow 2}}) = \sum_{i=1}^n der(V_i) \quad (8)$$

We apply this framework to analyze security across distinct bridge implementation layers, including i) on-chain contracts (handling token locking, minting, and withdrawal), ii) off-chain relayers (notary or validator networks responsible for cross-chain verification) and iii) finality mechanisms (ensuring transaction irreversibility and bridging safety). This layered security model enables targeted risk mitigation by isolating attack surfaces based on implementation constraints.

3.6 Attack Surfaces of Blockchain Bridges

Blockchain bridges rely on multiple components (on-chain smart contracts, off-chain relayers, and destination chain mechanisms), each introducing distinct security risks. We formalize these risks through attack surface analysis and provide a detailed summary of the identified vectors in the supplementary material.

3.6.1 Source Chain Layer Security. The attack surface of a bridge's source chain implementation is given by $\mathbb{I}_{src} = \langle \mathbb{T}_{src}, R_{src} \rangle$. For a bridge deploying n smart contracts on the source chain, we model its attack surface as $\Sigma_{src} = \langle \{SC_1, SC_2, \dots, SC_n\}, \mathbb{T}_{src} \rangle$ (i.e., contract implementations and trusted entities), with an attack surface area $Area(\Sigma_{src}) = \sum_{i=1}^n der(SC_i)$. For validator-controlled implementations ($SC_i = \emptyset$), the attack surface consists solely of the trust set $\mathbb{T}_{src}$ (i.e., the validators). While the contractual surface area is zero, the system remains exposed to validator-level risks such as economic incentives, collusion, or key compromise.

3.6.2 Off-Chain Layer Security. The off-chain layer, which handles bridge communication, consists of notaries, light clients, or sidechains, each with unique trust assumptions as we discuss next.

Notary-Based Mechanisms. For notary-based bridges ($\mathbb{T}_{off} = \text{notaries}$), security improves with more notaries, but decentralization is costly. If a sufficient fraction of notaries turn malicious, they can control the bridge for a duration $t > t^*$, enabling asset theft.

Light Client Mechanisms. For light-client bridges ($\mathbb{T}_{off} = \{LC, MP\}$), security depends on Merkle proofs (MP). If an attacker controls the light client, they can alter transaction verifications. The attack surface reduces to $Area(\Sigma_{off} \mid \mathbb{T}_{off} = \text{light client}) = der(\text{light client})$

Sidechain-Based Mechanisms. Sidechains use independent consensus rules (i.e., R_{src}). If R_{src} is compromised, an attacker can mint arbitrary tokens or modify the bridge state. If the sidechain guarantees R_{src} security, the attack surface vanishes $Area(\Sigma_{off} \mid \mathbb{T}_{off} = \text{sidechain}) = 0 \Leftrightarrow der(R_{src}) = 0$.

3.6.3 Destination Chain Layer Security. The destination chain's attack surface depends on its reliance on smart contracts (SC) and custodians $\Sigma_{dest} = \langle \tau_{dest}, \mathbb{T}_{dest} \rangle$, $\tau_{dest} = \{SC_1, SC_2, \dots, SC_n\} \cup \{\text{custodian}\}$. Attack surface area is $Area(\Sigma_{dest}) = n + 1$, assuming a custodian is present.

3.6.4 Total Attack Surface of a Blockchain Bridge. The total bridge attack surface is $\Sigma_{\mathbb{B}_{1 \leftrightarrow 2}} = \Sigma_{src} \cup \Sigma_{off} \cup \Sigma_{dest} \cup \Sigma_{other}$. Other attack vectors include updates to bridge protocols, governance failures, or rug-pulls. We validate our framework by classifying past major bridge exploits under this model.

Table 2: Categorized attack/disruption vectors across bridge layers. A dagger (†) marks vulnerabilities that only arise when a bridge architecture actually employs the relevant component (for example, an oracle).

Attack Vector	Source Chain	Off-Chain	Destination Chain
Contract Logic & Code Vulnerabilities			
V1: Reentrancy attacks	✓		✓
V2: Integer and arithmetic errors	✓		✓
V3: Access control and forged account flaws	✓		✓
V4: Race condition attacks	✓		✓
V5: Unsafe external call exploits	✓		✓
V6: Malicious event log manipulation	✓		✓
V7: Contract upgrade risks	✓		✓
Authentication & Authorization Failures			
V8: Fake burn/lock proofs	✓	✓	✓
V9: Malicious transaction modification	✓	✓	✓
V10: Light-client verification flaws	✓	✓	✓
V11: Oracle manipulation	✓†	✓	✓†
V12: Malicious custodian manipulation	✓†	✓	✓†
V13: Private key leakage or theft	✓†	✓†	✓†
Replay, Race, and Timing-Based Attacks			
V14: Timestamp manipulation	✓		✓
V15: Replay attacks		✓	✓
Consensus & Infrastructure Risks			
V16: Consensus failure (51% attack)	✓		✓
V17: Delayed finality exploitation	✓		✓
V18: Validator equivocation or misbehavior	✓†		✓†
V19: Denial of Service attacks	✓	✓	✓
V20: Deep chain reorganization	✓		✓
V21: Unbounded withdrawal limits	✓		✓
V22: Rugpull		✓	✓†
Front-End and Off-Chain Manipulation			
V23: Front-end deception		✓	

4 BRIDGE DESIGN PATTERNS AND IMPLEMENTATION LANDSCAPE

To contextualize our formalism and threat model, we survey representative bridge implementations, analyzing their design across three key layers: source chain, off-chain coordination, and destination chain. The supplementary material provides a detailed summary of their implementation patterns, trust assumptions, functional types, and blockchain coverage, along with full protocol descriptions and an overview of the broader ecosystem.

4.1 Static Analysis of Bridges

Table 3: Access control and code structure metrics of bridge smart contracts

Bridge Name	Local vars	Inheritances	Modifier Count	RoleBased	Standard Libs
Avalanche BridgeToken	2	1	0	Yes	1
Wormhole BridgeImplementation	58	2	1	No	3
Arbitrum L1GatewayRouter	9	5	2	No	0
Arbitrum L1ERC20Gateway	3	1	1	No	1
Arbitrum L1CustomGateway	5	2	2	No	1
Arbitrum L1WethGateway	0	1	0	No	2
Stargate Router	30	3	1	No	3
DeBridge DeBridgeGate	44	5	4	Yes	1
Across HubPool	50	5	3	No	3
Stargate Bridge	35	3	1	No	1
Allbridge Bridge	12	4	0	No	1
Nomad BridgeToken	6	4	0	No	0
Base L1StandardBridge	1	2	0	No	0
Optimism L1StandardBridge	0	2	0	Yes	0
Hyperliquid Bridge2	66	2	0	Yes	5
Meson BridgeV2	2	2	4	Yes	3

To assess the security robustness of bridge smart contracts, we conduct a *static analysis* across key dimensions drawn from our layered attack surface model. Static analysis focuses on examining the

Table 4: Lines of code (LOC), function visibility, and variable usage in bridge smart contracts

Bridge Name	LOC	Total lines	Public Funs	External Funs	Internal Funs	Private Funs	Global vars Declared
Avalanche BridgeToken	155	226	9	0	0	1	6
Wormhole BridgeImplementation	630	776	19	3	14	0	0
Arbitrum L1GatewayRouter	208	305	5	4	1	0	2
Arbitrum L1ERC20Gateway	112	161	5	2	0	0	6
Arbitrum L1CustomGateway	165	243	5	3	0	0	6
Arbitrum L1WethGateway	78	92	2	1	4	0	2
Stargate Router	280	323	0	17	4	0	5
DeBridge DeBridgeGate	905	1110	7	23	15	0	27
Across HubPool	610	1076	24	2	13	0	16
Stargate Bridge	249	310	2	14	4	0	9
Allbridge Bridge	215	320	1	11	2	0	5
Nomad BridgeToken	150	245	7	5	0	0	6
Base L1StandardBridge	260	321	8	8	6	0	2
Optimism L1StandardBridge	220	324	3	9	6	1	3
Hyperliquid Bridge2	570	840	17	15	2	14	18
Meson BridgeV2	167	210	8	8	0	2	4

Table 5: Analysis of external call behavior and defensive programming practices in bridge implementations

Bridge Name	Ext Funs	Low-level	Untrusted	Reentry Guard	Require/Assert	Custom Errors	Checks/Fn
Avalanche BridgeToken	5	0	0	No	14	0	1.56
Wormhole BridgeImplementation	2	5	0	Yes	21	0	0.58
Arbitrum L1GatewayRouter	3	0	0	No	9	0	1
Arbitrum L1ERC20Gateway	2	1	0	Yes	3	0	0.43
Arbitrum L1CustomGateway	4	0	0	Yes	6	0	0.75
Arbitrum L1WethGateway	0	0	0	No	3	0	0.43
Stargate Router	15	0	0	No	8	0	0.38
DeBridge DeBridgeGate	18	3	3	Yes	26	20	0.87
Across HubPool	13	3	3	Yes	26	0	1
Stargate Bridge	13	0	0	No	8	0	0.5
Allbridge Bridge	5	0	0	No	11	0	0.9
Nomad BridgeToken	4	0	0	No	4	0	0.33
Base L1StandardBridge	4	0	0	No	4	0	0.25
Optimism L1StandardBridge	4	0	0	No	0	0	0
Hyperliquid Bridge2	7	0	0	Yes	19	0	1
Meson BridgeV2	7	1	1	No	7	0	0.58

contract’s code without executing it, allowing us to detect vulnerabilities in logic, structure, and access control that may compromise the security priors *bridge causality*, *consistency*, and *token peg integrity*. We analyze the bridges in terms of i) access control and code structure metrics (Table 3), ii) function visibility and variable usage (Table 4) and iii) call behavior and defensive programming constructs (Table 5). This analysis maps to several high-impact attack vectors from Table 2, notably: V1: reentrancy attacks, V3: access control and forged account flaws, V5: unsafe external call exploits, V7: contract upgrade and misconfiguration risks. We identify several notable patterns and anomalies that offer insights into the heterogeneous security postures and implementation philosophies across bridges.

One immediate observation in Table 5 is the inconsistent use of reentrancy protection mechanisms. Despite having multiple externally callable functions, contracts such as *Stargate Router* (15 external functions) and *Stargate Bridge* (13 external functions) do not implement any form of reentrancy guard. In contrast, similarly sized contracts like *DeBridgeGate* and *Across HubPool* use such guards appropriately. The absence of reentrancy protection in these high-exposure contracts reflects a reliance on architectural assumptions.

We also find a curious disconnect in Table 3 between the use of role-based access control and the deployment of Solidity modifiers. For example, *Hyperliquid Bridge2* and *Avalanche BridgeToken* both report role-based access (“RoleBased = Yes”) but have zero modifiers. This suggests that access restrictions may be implemented inline via “require” statements rather than modularized through modifiers, leading to less auditable and reusable control logic. In contrast, *Nomad BridgeToken* and *Allbridge* show zero modifier usage and lack role-based protection, increasing susceptibility to unauthorized

access or logic misbehavior (V3). *Nomad*'s exploit history confirms this: a faulty initialization allowed any user to spoof legitimate senders.

A particularly unusual case in Table 5 is the *Optimism L1Standard Bridge* contract, which lacks any require or assert statements despite being externally callable and declaring role-based control. While this absence would typically raise concerns (since most bridges include basic sanity checks), it reflects a deliberate design choice. Optimism and *Base L1StandardBridge* rely on rollup-native architecture, where Ethereum layer-1 consensus enforces external validation. This structural safeguard reduces reliance on in-contract defensive coding against V3 and V5 vectors, but the lack of internal guards still poses risks if such contracts are reused outside this tightly scoped context.

In terms of architectural modularity, some contracts in Table 4 show a preference for deeply internalized logic structures. *Wormhole BridgeImplementation* and *Hyperliquid Bridge2* contain large numbers of internal or private functions (13-15 each), indicating encapsulated logic. While this may reflect thoughtful separation of concerns, it also complicates external audits and transparency.

We further observe disproportionate uses of access control infrastructure relative to contract size in Tables 3 and 4. *Meson BridgeV2*, with just 167 lines of code, defines four modifiers and a complete role-based access layer. In contrast, *Stargate Router* spans over 280 lines but employs one modifier. Such disparity shows divergent security philosophies among bridge developers, with some opting for granular control at all costs and others prioritizing operational simplicity.

A notable structural pattern in Table 4 appears in the allocation of global versus local variables. *DeBridgeGate* maintains a high number of global state variables (27), reflecting rich on-chain logic and persistent data tracking. Conversely, *Wormhole Bridge* holds 58 local variables but defines no global state, potentially due to reliance on proxy patterns or off-chain state.

Low-level calls (`call`, `delegatecall`, `staticcall`) introduce attack surface via vector V5. *Wormhole*, *DeBridge*, and *Across* use them multiple times, sometimes in the presence of untrusted inputs. While `call` can be necessary (e.g., for token forwarding), its misuse or failure to handle return values correctly has led to major incidents such as the *Poly Network* and *Qubit* [32] exploits. Bridges like *Avalanche BridgeToken*, *Nomad*, and *Allbridge* avoid low-level calls entirely, reducing exposure to unsafe execution paths.

Lastly, we compute the average number of checks per function (`#Checks/#Functions`) as a coarse proxy for defensive programming density. Contracts such as *Avalanche BridgeToken* (1.56) and *Hyperliquid Bridge2* (1.00) stand out as aggressively guarded, while *Optimism L1StandardBridge* (0.00) again reflects an absence of such measures. These findings indicate varying degrees of maturity in secure smart contract engineering practices.

Static analysis reveals that bridge contracts vary widely in their defensive quality. *DeBridgeGate* sets a high bar for smart contract hygiene, leveraging modifiers, role-based controls, and comprehensive assertions to guard its complex logic. *Wormhole* and *Across* balance complexity with modular defenses but remain susceptible to implementation flaws, as demonstrated in *Wormhole*'s historical signature validation failure. *Stargate*, *Nomad*, and *Allbridge* show critical gaps in protective constructs, leaving them exposed

to known attack vectors. These are particularly concerning given their public deployment.

4.2 Transaction Analysis of Bridges




Figure 3: Attack flow of the Ronin Bridge exploit. The attacker gained control of 5 out of 9 Ronin Bridge validators and used them to sign and submit forged withdrawals. The stolen funds amounted to 173,600 ETH, along with an additional 25.5 million USDC that was exchanged for 8,564 ETH.

To better understand the on-chain behavior of bridge exploiters, we perform a transaction-level analysis focused on Ethereum addresses used by attackers, as Ethereum is commonly chosen for fund collection due to its exchange access and availability of mixing services. We obtain Ethereum transaction data by running a full node using Geth (<https://github.com/ethereum/go-ethereum>) and parsing all transactions via Ethereum-ETL (<https://github.com/blockchain-etl/ethereum-etl>). We visualize the attack patterns in transaction graphs to capture attacker behavior. The analysis includes both externally owned account transactions and internal transactions, covering pre-attack, attack, and post-attack phases, where the phase lengths span the full Ethereum history up to March 2025.

Due to space limitations, an example transaction subgraph is shown in Figure 3, and seven other notable attacks are visualized in ?? In the figure, the *Ronin bridge* to Ethereum was victim to a validator key compromise in an attack now attributed to North Korea's Lazarus Group. *Ronin* used a 5-of-9 multisig validation for bridge withdrawals. Attackers compromised five private keys through social engineering. Once in control, they issued two fraudulent transactions, draining 173,600 ETH and 25.5M USDC (worth approximately \$624 million at the time) from the *Ronin bridge* in a single stroke. This event, the largest DeFi hack ever, was essentially a failure of the bridge's trust model: the off-chain validators were assumed honest, but the minimal quorum and centralized key management (a single entity controlled 4 of 9 validators) made it easy for an attacker to breach. The breach went unnoticed for six days until a user discovered it.

Most exploiter addresses have little to no previous on-chain activity. In many cases, the wallets were newly created and funded with just a sufficient amount to cover gas fees. The average number of pre-attack transactions was 28, suggesting that even the more

“active” wallets were only lightly used. Some of these initial funds even originated from mixers, most notably Tornado Cash [33] in the *Wormhole* bridge hack, where the exploiter received ETH from Tornado Cash before launching the attack, an indication of deliberate obfuscation of the funding source. In other instances, such as the *Nomad* and *Horizon bridge* [34] hacks, attackers used unlabeled intermediate wallets, further complicating traceability, and as a strategy to hide the funding sources.

During the attack period, the behavior is highly focused. The attacker often interacts directly with the bridge contract via a few function calls, most of the time, only one forged withdrawal or proof verification. In attacks such as *Nomad*, we observe many small transactions from multiple addresses exploiting the same vulnerability after it became public, showing replication once the vulnerability was made public.

Internal transactions are important for creating the contract call sequences. The contract call traces highlight how the attack is carried out at the smart contract level, often using call stacks involving specific functions like *withdrawERC20For()* or *verifyProof()*. This helps verify the exploit strategies and confirms the role of compromises or logic flaws.

After the exploit, the stolen funds are moved rapidly to avoid detection. Some exploiters transferred stolen funds through multiple intermediate wallets before returning them to a central exploiter address, for instance, in the *Ronin Bridge* hack. Some others bridged funds to other blockchains such as Ethereum, Polygon, and Avalanche using bridges. Many proceeds are laundered through privacy tools like Tornado Cash or converted via decentralized exchanges (DEXs).

After initial laundering, several attacker addresses remain active. In our analysis, 6 out of 18 wallets continued transacting for more than one year after the exploit, indicating possible long-term usage or reactivation of compromised wallets.

Although certain early detection remains challenging, our analysis identifies behavioral clues that may serve as early warning signs of an upcoming exploit. A common pattern is the initial funding of a wallet with just enough ETH to cover gas fees, which is sourced from mixers. These wallets usually indicate little to no prior activity. Moreover, in order to verify system reactions or contract behavior, numerous low-value probing transactions are occasionally issued prior to the main exploit. These patterns could be leveraged to flag suspicious activity in near-real time.

5 SECURITY ANALYSIS OF BRIDGE ATTACKS

Our analysis reveals that bridge causality and consistency priors are violated in bridge attacks. Most attacks involve violations of the cross-chain causality prior. While Equation (4), the peg prior, is formally violated as a result of these attacks, the peg itself is not the target of the attacks. That is, attackers do not manipulate token prices or attempt to create price divergence across chains. Instead, the violations arise indirectly because tokens are minted or released without proper backing, breaking the assumption of value parity that the peg prior encodes. Due to space limitations, detailed attack descriptions and a meta-analysis are provided in the supplementary material.

Attack Vector Analysis. ?? shows that V3 (access control) and V13 (key leakage) have been exploited 10 times each. The exploits reveal two broad categories of failures: (1) Off-chain trust failures and (2) On-chain validation failures. The off-chain trust failures encompass all incidents where the bridge’s security relied on individuals or off-chain systems that were compromised, notably the multisig key compromises (*Ronin*, *Harmony*, *Multichain*, *Orbit*) and related cases. In *Ronin* and *Harmony*, the sources of failure were outside the blockchain: hackers penetrated the organizations controlling the validators and obtained the private keys needed to sign fake transactions. These attacks did not exploit a bug in code; they exploited insufficient decentralization and operational security. Essentially, the assumption that a small set of validators would remain honest was violated. When 5 of 9 *Ronin* validators and 2 of 5 *Harmony* signers turned malicious (via key theft), the bridge smart contracts on-chain duly obeyed the malicious signatures, an example of how improperly validated inputs led the system to execute unintended behavior. Similarly, *Multichain*’s collapse was an off-chain failure: the system’s architecture secretly concentrated too much trust in one individual. These illustrate that trusted or trust-minimized bridges are vulnerable by design. They introduce new trust points (keys, signers, servers) that attackers target through phishing, insider collusion, malware and more. The on-chain validation failures, on the other hand, cover exploits like *Poly Network*, *Wormhole*, *Nomad*, *BSC Token Hub*, *Qubit*, where the bridge smart contracts or crypto verification routines on one of the chains had a flaw. In these cases, the attacker manipulated the code logic (e.g., forging a message that the contract erroneously accepted as valid). For *Poly*, the bug was an unchecked external call in a privileged contract; for *Wormhole*, a bypassed signature verification on Solana; for *Nomad*, an incorrect initialization setting that trusted everyone; for BSC, a flawed light-client proof verification.

Despite different mechanics, these represent software/security bugs in the bridge implementation. The failures occurred either on the source chain contract (for *Poly* on Ethereum) or the destination chain contract (for *Wormhole*, *Nomad*, *Qubit*) or the intermediate relay logic (as in BSC’s light client). Crucially, these attack vectors map to points of validation in the bridge architecture.

In several cases (*Wormhole*, *Nomad*, *BSC*), the violated core assumption was that the smart contract correctly validates the cross-chain proof. These were not fundamental cryptographic failures (algorithms such as *secp256k1* were sound), but errors in how the algorithms were applied. This suggests many bridge exploits are avoidable with more rigorous software engineering, such as comprehensive audits, formal verification of bridge contracts, and in-depth defense (e.g., requiring multiple independent checks of a proof).

Importantly, many of these incidents exhibit high damage-to-effort ratios. The impact of a successful exploit has often reached hundreds of millions of dollars, while the effort required, whether key theft, phishing, or exploiting a poorly audited contract, remains relatively low. These high $\text{der}(V)$ values underscore why such attack vectors persist across different bridge designs and why they dominate the historical incident landscape.

RQ1: Are bridge attacks structurally preventable, or can they be mitigated by design? Yes, the evidence suggests that certain designs are far more robust. Notably, **no major exploits have occurred**

on fully trustless light-client bridges or rollup bridges. For example, Cosmos’s L1-to-L1 IBC channels have operated without incident between dozens of chains since 2021 (a lossless reentrancy bug was discovered early [35]), and Ethereum’s (young) rollup L1-to-L2 bridges have not been breached to date. Avalanche L1-to-L1 bridge to Ethereum continues functioning without any incidents. These systems benefit from minimal attack surface: there are no external signers to target, and the verification logic is often simpler or inherited from consensus (IBC uses tendermint-like light clients, which are well-vetted; rollups use post validity/fraud proofs checked by L1). This suggests that trustless designs inherently eliminate entire classes of attacks (notably, attack vector V3 of Table 2). These architectures inherently exclude common attack vectors like compromised keys or poorly controlled access mechanisms, offering significantly lower $\text{der}(V)$ values.

By contrast, most attacks struck bridges that added a new layer of validation on top of the two blockchains. Each additional component (e.g., a multisig, an off-chain oracle network, or a novel smart contract) introduced complexity and potential weakness. Our analysis shows evidence that the more a bridge can lean on intrinsic blockchain security (native verification), the safer it will be. In contrast, trust-minimized L1-to-L1 bridges are often only as strong as their validator set, which could be much weaker than either chain’s consensus. A quantified framework by LIFI found in 2022 that many so-called trust-minimized bridges still effectively rely on a majority of a few validators, leaving them vulnerable to 51% attacks or key compromises [36]. Indeed, *Ronin* and *Harmony* hacks demonstrated that a <10-member validator set offers poor resilience. Increasing the set (e.g., Axelar has 70+ validators) improves security, but unless those validators are as robust as a full blockchain consensus (with hundreds of nodes and economic security), the bridge remains a tempting honeypot for hackers. Therefore, trust-minimized bridges offer better resilience than naive trusted models, but they still concentrate trust among a group.

Properly decentralized networks with bonding (to slash malicious validators) are harder to attack; for instance, a hack of the scale of *Ronin* on a larger bonded validator bridge has not occurred yet. However, Ethereum-to-Solana (L1-to-L1) *Wormhole*’s Feb 2022 case reminds us that even absent collusion, a bug in the validation code can be just as catastrophic. This blurs the line between *fundamental* and *implementation* issues: we argue that the complexity of blockchain bridge protocols is a fundamental weakness, as more complex logic yields a higher chance of errors.

Finally, the prevalence of hacks has spurred research into failure-resistant bridges. For example, designs like “circuit breaker” bridges propose that if an unusually large withdrawal is attempted, the bridge only allows it after a delay or community vote, to mitigate instant draining. Others suggest blending multiple validation mechanisms, e.g., requiring both an off-chain multi-sig and an on-chain light client to agree, which could guard against either one failing alone (at the cost of added complexity). From a benchmarking perspective, we argue that diversity in validation (multi-layer security) might drastically reduce the probability of a successful exploit, albeit with performance trade-offs. Moreover, as bridges adopt techniques like zero-knowledge (zk) proofs for validation (e.g., *zkBridge* frameworks), some traditional attack vectors might be closed (a zk proof can attest to a source chain state without any trusted signer

at all). Such bridges could offer the holy grail: trustless interoperability between any two chains, with cryptographic guarantees and no reliance on human validators. Early prototypes (like *zkRelay* and light-client circuits) are promising, but they require heavy computation and are just entering practical deployment.

RQ2: Are bridge attacks detectable in real-time? A surprising inability of the blockchain ecosystem is the lack of a strong analytics foundation that could monitor bridges for potential attacks. In the *Ronin* hack, it took days (and a user complaint) to notice that \$620M had vanished. An obvious insight here is that real-time auditing and alerts are essential, but it raises a more fundamental question: why have such systems not been developed, despite the substantial capital invested in decentralized finance?

We outline four reasons. First, many DeFi projects operate with lean teams focused primarily on product development and protocol maintenance, with their expertise originating from smart contract development. Allocating resources to build and maintain sophisticated monitoring systems often falls outside their immediate capabilities. Second, implementing real-time detection systems requires a scalable analytics infrastructure capable of handling large volumes of blockchain data. This necessitates expertise in data engineering and analytics, skills that may not be prevalent among smart contract developers. Third, analyzing blockchain data presents unique challenges, including model scalability and the accuracy of anomaly detection. Traditional data analytics methods are often ill-suited to address the decentralized and heterogeneous structure of blockchain networks. Moreover, only a limited number of research efforts have focused on foundational questions specific to blockchain, such as identifying influential addresses within transaction networks (e.g., Alphacore [37]; see [38] for a recent survey). Lastly, investing in detection systems does not directly generate revenue, making it less appealing for projects to allocate funds toward such initiatives. While white-hat hacking offers some incentives, it often involves high effort with uncertain rewards. Nonetheless, advancements in automated security testing, particularly fuzzing [39], are emerging as promising solutions. Fuzzing involves providing invalid, unexpected, or random data inputs to smart contracts to uncover vulnerabilities.

The anomaly detection effort must eventually fall on bridge managers. Some bridges (e.g., *Gravity*, *Chainlink CCIP* [40]) are now deploying independent, on-chain monitoring tools that can pause a bridge if suspicious activity is detected (for example, if an invariant is violated or if an unrecognized validator signature appears). In centralized finance, analogous systems flag large transfers for manual review; bridges could implement similar safeguards or circuit-breakers, but this reintroduces trust if not done in a decentralized way.

RQ3: What mechanisms exist to mitigate damage and minimize asset loss? There is no standard set of mechanisms for limiting damage once a bridge is under attack, and mitigation strategies remain ad hoc and uneven across protocols. *Gravity* and *Chainlink* are developing anomaly detection tools that can halt bridge operations in the face of an attack. In *Celer*, there is a buffer delay period during which transactions can be independently verified before a transfer is finalized. If inconsistencies are detected, the transfer can be halted, offering an added layer of fraud prevention. Post-factum, blockchain

transparency does aid forensics; after these hacks, investigators (e.g., *ZachXBT* on Twitter) traced funds through addresses and often identified suspects or recovered portions. But prevention and rapid response are clearly preferable to relying on clawbacks. Benchmark evaluations must incorporate attack detection latency as a metric; i.e., how quickly a breach is noticed and halted. Shorter detection times (as in BSC’s case, hours) can significantly limit losses compared to delayed discovery (as in *Ronin*).

The bridge attacks of 2021–2025 underscore a core tension in blockchain interoperability: security vs. connectivity. Bridges expand what users can do (moving assets across chains), but they also expand the attack surface beyond any single blockchain. Our survey of exploits indicates that while many were due to low-hanging fruit bugs or a lack of operational security, there are also intrinsic risks whenever custody of funds shifts to a separate system.

RQ4: What is the long-term outlook for bridge security? What foundational steps are needed to design formally verifiable or provably secure bridge protocols? If truly trustless (or provably secure) bridge mechanisms can be standardized, we expect far fewer catastrophic failures. Until then, however, bridge designers must assume that attacks are a matter of when, not if, and design accordingly. This means employing rigorous audits, fail-safes, decentralized validation, and continuous monitoring as part of any cross-chain system. Past events suggest that increasing decentralization and using established, simple verification methods (light clients, chain consensus) correlate with higher security, whereas highly complex or centralized bridges have correlated with the largest failures.

On the other hand, the multisig key hacks raise the question of whether those bridges were destined to fail, i.e., was it an inevitability that eventually some signer’s key would get compromised? One can argue that any valuable honeypot (whether a contract with billions of locked value or a key controlling billions) will, over time, face relentless attack until a weakness is found. Even with multi-party validators, if the set is small or the keys are not protected by robust hardware security modules and operational security, attackers have a focal point to strike. Thus, fundamental architecture choices (like using a 2-of-5 multisig with hot keys) can be seen as structural vulnerabilities, not in the theoretical sense, but in practical terms of presenting an irresistibly weak link (humans with keys) compared to the surrounding blockchain’s security.

A lesson learned with pain is that trusted bridges must be avoided due to significant management risks [2]. Rollups offer a trustless solution but are limited to chains with compatible runtimes. New protocols like the *IBC* protocol of Cosmos and *Chainlink*’s Cross-Chain Interoperability Protocol aim to bridge heterogeneous chains. However, even those new-generation solutions introduce varying trust assumptions. IBC uses light clients when both chains natively support it. However, when adapted for non-IBC-compatible chains, additional trust assumptions are required. CCIP, on the other hand, relies on Chainlink’s decentralized oracle network to facilitate cross-chain interactions, introducing a trust-minimized model where users depend on the honesty and reliability of the oracle nodes.

6 CONCLUSION

Blockchain bridges have become indispensable infrastructure for cross-chain interoperability, yet they remain among the most vulnerable components of the decentralized ecosystem. Our study presents the first comprehensive, data-driven systematization of bridge security, combining formal modeling, large-scale static analysis, and empirical transaction-level investigations. We show that the majority of bridge attacks violate core security priors, particularly causality and consistency, without breaching the value peg itself.

Our analysis reveals that most successful exploits stem from two dominant classes: off-chain trust failures and on-chain validation bugs. These vectors frequently exhibit high damage-to-effort ratios, which we formalize through the $\text{der}(V)$ metric. We also find that architectural design plays a decisive role in resilience. Trustless models, such as light-client and rollup-native bridges, have so far withstood real-world adversarial conditions, while trusted and loosely trust-minimized bridges remain disproportionately vulnerable. Emerging defense mechanisms like circuit breakers, buffer delays, and hybrid validation schemes offer promise but are inconsistently implemented and lack formal standardization.

Finally, we identify critical gaps in real-time detection and mitigation. Despite billions in locked value, most bridges lack robust monitoring infrastructure or containment protocols, and responses to attacks are often delayed and improvised. Bridging this gap will require new research into on-chain anomaly detection, decentralized fail-safe mechanisms, and benchmarking frameworks that account for detection latency and systemic risk.

We hope this work provides a foundation for rethinking how cross-chain systems are built, evaluated, and secured. As bridges continue to evolve, their long-term viability will depend not only on throughput and composability but also on their ability to withstand failure in adversarial conditions. Bridging across blockchains should not mean bridging across trust assumptions.

REFERENCES

- [1] “Bridge Volume - DefiLlama.” <https://defillama.com/bridges>. Accessed: 2025-05-29.
- [2] B. Kiepuszewski, “Rollups are the most secure bridges.” <https://archive.devcon.org/devcon-6/rollups-are-the-most-secure-bridges/>, 2022. Devcon 6, Bogota, October 13, 2022.
- [3] “Multichain.” <https://multichain.org/>. Accessed: 2025-04-24.
- [4] “Ronin bridge.” <https://docs.roninchain.com/apps/ronin-bridge>, 2025. Accessed: 2025-05-21. Official documentation for the Ronin Bridge, facilitating ERC-20 token and NFT transfers between Ethereum and the Ronin chain.
- [5] “Rainbow bridge.” <https://rainbowbridge.app/>, 2025. Accessed: 2025-05-21. Official interface for the Rainbow Bridge, facilitating trustless and permissionless asset transfers between Ethereum, NEAR, and Aurora blockchains.
- [6] “Polygon plasma bridge.” <https://docs.polygon.technology/>, 2025. Accessed: 2025-05-21. Official documentation for the Polygon Plasma Bridge, enabling secure asset transfers between Ethereum and Polygon using Ethereum’s Plasma scaling framework.
- [7] A. Augusto, R. Belchior, J. Pfannschmidt, A. Vasconcelos, and M. Correia, “Xchain-watcher: Monitoring and identifying attacks in cross-chain bridges,” *arXiv preprint arXiv:2410.02029*, 2024.
- [8] J. Wu, K. Lin, D. Lin, B. Zhang, Z. Wu, and J. Su, “Safeguarding blockchain ecosystem: Understanding and detecting attack transactions on cross-chain bridges,” *arXiv preprint arXiv:2410.14493*, 2024.
- [9] N. Belenkov, V. Callens, A. Murashkin, K. Bak, M. Derka, J. Gorzny, and S.-S. Lee, “Sok: A review of cross-chain bridge hacks in 2023,” *arXiv preprint arXiv:2501.03423*, 2025.
- [10] N. Li, M. Qi, Z. Xu, X. Zhu, W. Zhou, S. Wen, and Y. Xiang, “Blockchain cross-chain bridge security: Challenges, solutions, and future outlook,” *Distributed Ledger Technologies: Research and Practice*, vol. 4, no. 1, pp. 1–34, 2025.

- [11] "Ethereum." <https://ethereum.org/>, 2025. Accessed: 2025-05-21. Official website of Ethereum, a decentralized platform supporting smart contracts and decentralized applications (dApps).
- [12] C. Huang, T. Yan, and C. J. Tessone, "Seamlessly transferring assets through layer-0 bridges: An empirical analysis of stargate bridge's architecture and dynamics," in *Companion Proceedings of the ACM Web Conference 2024*, pp. 1776–1784, 2024.
- [13] "Polygon." <https://polygon.technology/>, 2025. Accessed: 2025-05-21. Official website of Polygon, an Ethereum scaling platform offering Layer 2 solutions such as the PoS chain, zkEVM, and other modular frameworks for building interoperable blockchain networks.
- [14] Binance, "Binance smart chain whitepaper," 2020.
- [15] "Avalanche." <https://avax.network/>, 2025. Accessed: 2025-05-21. Official website of Avalanche, a high-performance Layer 1 blockchain platform developed by Ava Labs, featuring a unique three-chain architecture (X-Chain, C-Chain, and P-Chain) and the Avalanche Consensus Protocol for rapid transaction finality.
- [16] "Arbitrum." <https://arbitrum.io/>, 2025. Accessed: 2025-05-21. Official website of Arbitrum, a Layer 2 scaling solution for Ethereum utilizing optimistic rollups to enhance transaction throughput and reduce fees while maintaining security.
- [17] "Optimism." <https://www.optimism.io/>, 2025. Accessed: 2025-05-21. Official website of Optimism, a Layer 2 scaling solution for Ethereum utilizing optimistic rollups to enhance transaction throughput and reduce fees while maintaining security.
- [18] S. Subramanian, A. Augusto, R. Belchior, A. Vasconcelos, and M. Correia, "Benchmarking blockchain bridge aggregators," in *2024 IEEE International Conference on Blockchain (Blockchain)*, pp. 37–45, IEEE, 2024.
- [19] A. Augusto, A. Vasconcelos, M. Correia, and L. Zhang, "Xchaindatagen: A cross-chain dataset generation framework," *arXiv preprint arXiv:2503.13637*, 2025.
- [20] A. Augusto, R. Belchior, M. Correia, A. Vasconcelos, L. Zhang, and T. Hardjono, "Sok: Security and privacy of blockchain interoperability," in *2024 IEEE Symposium on Security and Privacy (SP)*, pp. 3840–3865, IEEE, 2024.
- [21] "Nomad." <https://app.nomad.xyz/>, 2025. Accessed: 2025-05-21. Official website of Nomad, a security-first cross-chain messaging protocol enabling secure and gas-efficient communication between blockchains.
- [22] D. Lin, Z. Zheng, J. Wu, J. Yang, K. Lin, H. Xiao, B. Song, and Z. Zheng, "Track and trace: Automatically uncovering cross-chain transactions in the multi-blockchain ecosystems," *arXiv preprint arXiv:2504.01822*, 2025.
- [23] "Bitcoin." <https://bitcoin.org/>, 2025. Accessed: 2025-05-21. Official website of Bitcoin, the first decentralized cryptocurrency and blockchain network.
- [24] "Solana." <https://solana.com/>, 2025. Accessed: 2025-05-21. Official website of Solana, a high-performance blockchain platform supporting decentralized applications and smart contracts.
- [25] "Orbs network." <https://www.orbs.com/>, 2025. Accessed: 2025-05-21. Official website of Orbs Network, a decentralized Layer-3 blockchain infrastructure designed to enhance the capabilities of existing smart contracts by providing decentralized backend services.
- [26] "Xai games." <https://xai.games/>, 2025. Accessed: 2025-05-21. Official website of Xai Games, a Layer 3 blockchain platform built on Arbitrum, designed to bring AAA gaming experiences to Web3 by enabling seamless integration of blockchain technology into games.
- [27] "zksync." <https://zksync.io/>, 2025. Accessed: 2025-05-21. Official website of zkSync, a Layer 2 Ethereum scaling solution utilizing zero-knowledge rollups to enhance scalability and reduce transaction costs while maintaining high security.
- [28] "Wormhole." <https://wormhole.com/>, Accessed: 2025-04-24.
- [29] "Loopring." <https://loopring.org/>, 2025. Accessed: 2025-05-21. Official website of Loopring, a Layer 2 zkRollup protocol on Ethereum that enables scalable, secure, and low-cost decentralized trading and payments.
- [30] Connex, "Connex: The interoperability protocol for cross-chain applications." <https://www.connex.network/>, 2025. Accessed: 2025-06-03.
- [31] Across, "Across: Cross-chain bridge for ethereum and layer 2 networks." <https://across.to/>, 2025. Accessed: 2025-06-03.
- [32] "Qubit finance." <https://qbt.fi/>, 2025. Accessed: 2025-05-21. Official website of Qubit Finance, a decentralized money market platform on Binance Smart Chain (BSC) that connects lenders and borrowers efficiently and securely.
- [33] "Tornado cash." <https://tornado.cash/>, 2025. Accessed: 2025-05-21. Official website of Tornado Cash, a decentralized, non-custodial privacy protocol on Ethereum that uses zero-knowledge proofs to enable anonymous transactions.
- [34] "Horizon bridge." <https://bridge.harmony.one/>, 2025. Accessed: 2025-05-21. Official website of Horizon Bridge, a cross-chain bridge developed by Harmony that facilitates the transfer of crypto assets between Ethereum, Binance Smart Chain, and Harmony networks.
- [35] B. Lindrea, "Cosmos patches 'critical' ibc protocol bug, saving \$126m," 2024. Accessed: 2025-05-19.
- [36] A. Chand and M. Murdock, "Li.fi: With bridges, trust is a spectrum," July 2022. LI.FI Knowledge Hub.
- [37] F. Victor, C. G. Akcora, Y. R. Gel, and M. Kantarcioglu, "Alphacore: data depth based core decomposition," in *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery & Data Mining*, pp. 1625–1633, 2021.
- [38] P. Azad, C. Akcora, and A. Khan, "Machine learning for blockchain data analysis: Progress and opportunities," *Distributed Ledger Technologies: Research and Practice*, 2024.
- [39] M. Rodler, D. Paaßen, W. Li, L. Bernhard, T. Holz, G. Karame, and L. Davi, "Ef/cf: High performance smart contract fuzzing for exploit generation," in *8th IEEE European Symposium on Security and Privacy, EuroS&P 2023, Delft, Netherlands, July 3–7, 2023*, pp. 449–471, IEEE, 2023.
- [40] Chainlink Labs, "Chainlink cross-chain interoperability protocol (ccip)." <https://chain.link/cross-chain>, 2025. Accessed: 2025-04-24.